

NOTE 1. CROSS ENTROPY LOSS

class torch.nn.CrossEntropyLoss(weight=None, size_average=True, ignore_index=-100, reduce=True)

This criterion combines *LogSoftMax* and *NLLLoss* is one single class. It is useful when training a classification problem with C classes. If provided, the optional argument *weight* should be a *1D Tensor* assigning weight to each of the classes. This is particularly useful when you have an unbalanced training set.

The *input* is expected to contain scores for each class. *input* has to be a *2D Tensor* of size $(minibatch, C)$. This criterion expects a class index (0 to $C-1$) as the *target* for each value of a *1D Tensor* of size of *minibatch*.

The loss can be described as:

$$loss(x, class) = -\log \frac{e^{x[class]}}{\sum_j e^{x[j]}} = -x[class] + \log \sum_j e^{x[j]}$$

or in the case of the *weight* argument being specified:

$$loss(x, class) = weight[class] \cdot (-x[class] + \log \sum_j e^{x[j]})$$

The losses are averaged across observations for each minibatch.

Parameter:

- *weight*(Tensor, optional) - a manual rescaling weight given to each class. If given, has to be a *Tensor* of size C
- *size_average*(bool, optional) - By default, the losses are averaged over observations for each minibatch. However, if the field *size_average* is set to *False*, the losses are instead summed for each minibatch. Ignored if *reduce* is *False*
- *ignore_index*(int, optional) - Specifies a target value that is ignored and doesn't contribute to the input gradient. When *size_average* is *True*, the loss is averaged over non-ignored targets
- *reduce*(bool, optional) - By default, the losses are averaged or summed over observations for each minibatch depending on *size_average*. When *reduce* is *False*, returns a loss per batch element instead and ignores

size_average. Default, *True*

Shape:

- Input: (N, C) where $C = \text{number of classes}$
- Target: (N) where each value is $0 \leq \text{targets}[i] \leq C - 1$
- Output: scalar. If reduce is *False*, then (N) instead

Example:

```
1 import torch
2 import torch.nn as nn
3
4 # y_pred.shape = torch.Size([2, 4])
5 # y.shape = torch.Size([2])
6 y_pred = torch.tensor([[0.1, 0.1, 0.2, 0.6], [0.1, 0.1, 0.2, 0.6]])
7 y = torch.tensor([1, 1])
8 print('y_pred: ', y_pred)
9 print('y: ', y)
10 criterion = nn.CrossEntropyLoss()
11 loss = criterion(y_pred, y)
12 print('loss: ', loss)
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL JUPYTER

```
PS D:\Downloads\QA-20220824T162725Z-001\QA> python -u "d:\Downloads\QA-20220824T1
y_pred: tensor([[0.1000, 0.1000, 0.2000, 0.6000],
               [0.1000, 0.1000, 0.2000, 0.6000]])
y: tensor([1, 1])
loss: tensor(1.5590)
```

$$\text{loss}(x_1, \text{class}_1) = -\log \frac{e^{0.1}}{e^{0.1} + e^{0.1} + e^{0.2} + e^{0.6}} = 1.5590$$

$$\text{loss}(x_2, \text{class}_2) = -\log \frac{e^{0.1}}{e^{0.1} + e^{0.1} + e^{0.2} + e^{0.6}} = 1.5590$$

$$\text{loss}(x, \text{class}) = \frac{1}{2} \cdot (\text{loss}(x_1, \text{class}_1) + \text{loss}(x_2, \text{class}_2)) = 1.5590$$